

# Multiple Linear Regression

Dr Liwan Liyanage - Western Sydney University

2 July 2018

# Multiple Linear Regression

First upload the data file

```
library(MASS)  
attach(Boston)
```

## Multiple Linear Regression

In order to fit a multiple linear regression model using least squares, we again use the `lm()` function. The syntax `lm(y~X1+X2+X3)` is used to fit a model with three predictors, X1, X2, and X3.

```
lm.fit2 = lm(formula = medv ~ lstat + age , data = Boston )  
lm.fit2
```

```
##
```

```
## Call:
```

```
## lm(formula = medv ~ lstat + age, data = Boston)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)          lstat          age
```

```
##      33.22276      -1.03207       0.03454
```

## Summary

```
summary(lm.fit2)
```

```
##
```

```
## Call:
```

```
## lm(formula = medv ~ lstat + age, data = Boston)
```

```
##
```

```
## Residuals:
```

##	Min	1Q	Median	3Q	Max
##	-15.981	-3.978	-1.283	1.968	23.158

```
##
```

```
## Coefficients:
```

##		Estimate	Std. Error	t value	Pr(> t )
##	(Intercept)	33.22276	0.73085	45.458	< 2e-16 ***
##	lstat	-1.03207	0.04819	-21.416	< 2e-16 ***
##	age	0.03454	0.01223	2.826	0.00491 **

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
```

```
##
```

# Anova

```
anova(lm.fit2)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: medv
```

##	Df	Sum Sq	Mean Sq	F value	Pr(>F)		
## lstat	1	23243.9	23243.9	609.955	< 2.2e-16 ***		
## age	1	304.3	304.3	7.984	0.004907 **		
## Residuals	503	19168.1	38.1				
## ---							
## Signif. codes:	0	'***'	0.001	'**'	0.01	'*' 0.05	'.' 0.1

# The Boston data set contains 13 variables,

First check if all variables are numeric?

```
sapply(Boston,class)
```

```
##      crim      zn      indus      chas      nox  
## "numeric" "numeric" "numeric" "integer" "numeric" "numer  
##      dis      rad      tax      ptratio      black      ls  
## "numeric" "integer" "numeric" "numeric" "numeric" "numer
```

## Note

It would be cumbersome to have to type all 13 variables in order to perform a regression using all of the predictors. Instead, we can use the following

```
lm.fit3 = lm(formula = medv ~ ., data = Boston )  
lm.fit3
```

```
##
```

```
## Call:
```

```
## lm(formula = medv ~ ., data = Boston)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)          crim              zn          indus  
##  3.646e+01   -1.080e-01   4.642e-02   2.056e-02   2  
##          nox          rm          age          dis  
## -1.777e+01   3.810e+00   6.922e-04  -1.476e+00   3  
##          tax      ptratio          black          lstat  
## -1.233e-02  -9.527e-01   9.312e-03  -5.248e-01
```

## Summary

```
summary (lm.fit3)
```

```
##
```

```
## Call:
```

```
## lm(formula = medv ~ ., data = Boston)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

##	-15.595	-2.730	-0.518	1.777	26.199
----	---------	--------	--------	-------	--------

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
```

## (Intercept)	3.646e+01	5.103e+00	7.144	3.28e-12	***
## crim	-1.080e-01	3.286e-02	-3.287	0.001087	**
## zn	4.642e-02	1.373e-02	3.382	0.000778	***
## indus	2.056e-02	6.150e-02	0.334	0.738288	
## chas	2.687e+00	8.616e-01	3.118	0.001925	**
## nox	-1.777e+01	3.820e+00	-4.651	4.25e-06	***



## We can access the individual components of a summary object by name

(type `?summary.lm` to see what is available).

- ▶ Hence `summary(lm.fit)$r.sq` gives us the  $R^2$ , and
- ▶ `summary(lm.fit)$sigma` gives us the RSE.

```
?summary.lm
```

```
## starting httpd help server ... done
```

```
summary(lm.fit3)$r.sq
```

```
## [1] 0.7406427
```

```
summary(lm.fit3)$sigma
```

```
## [1] 4.745298
```

## Variance Inflation Factors and Multicollinearity

The `vif()` function, part of the `car` package - Variance Inflation Factors can be used to compute variance inflation factors. Most VIF's are low to moderate for this data. Larger the VIF greater the multicollinearity.

The `car` package is not part of the base R installation so it must be downloaded the first time you use it via the `install.packages` option in R.

```
library(carData)
library(car)
vif(lm.fit3)
```

```
##      crim      zn    indus    chas      nox      rm
## 1.792192 2.298758 3.991596 1.073995 4.393720 1.933744 3
##      rad      tax  ptratio    black    lstat
## 7.484496 9.008554 1.799084 1.348521 2.941491
```

# Interpreting the Variance Inflation Factor

Variance inflation factors range from 1 upwards. The numerical value for VIF tells you (in decimal form) what percentage the variance is inflated for each coefficient. For example, a VIF of 1.9 tells you that the variance of a particular coefficient is 90% bigger than what you would expect if there was no multicollinearity

If there was no correlation with other predictors. A rule of thumb for interpreting the variance inflation factor:

- ▶ 1 = not correlated.
- ▶ Between 1 and 5 = moderately correlated.
- ▶ Greater than 5 = highly correlated.

$VIF = 1/(1 - R_i^2)$  where  $R_i^2$  is the  $R^2$  value when  $i$ th predictor is regressed against other predictors.

## What if we would like to perform a regression using all of the variables but one?

For example, in the above regression output, age has a high p-value. So we may wish to run a regression excluding this predictor. The following syntax results in a regression using all predictors except age.

```
lm.fit4=lm(medv~.-age ,data=Boston )  
summary(lm.fit4)
```

```
##
```

```
## Call:
```

```
## lm(formula = medv ~ . - age, data = Boston)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max  
## -15.6054  -2.7313  -0.5188   1.7601  26.2243
```

```
##
```

```
## Coefficients:
```

Alternatively, the `update()` function can be used.

```
lm.fit5=update(lm.fit3,~.-age)
summary(lm.fit5)
```

```
##
## Call:
## lm(formula = medv ~ crim + zn + indus + chas + nox + rm
##      rad + tax + ptratio + black + lstat, data = Boston)
##
## Residuals:
##      Min        1Q    Median        3Q       Max
## -15.6054  -2.7313  -0.5188   1.7601  26.2243
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  36.436927   5.080119   7.172 2.72e-12 ***
## crim        -0.108006   0.032832  -3.290 0.001075 **
## zn           0.046334   0.013613   3.404 0.000719 ***
## indus        0.020562   0.061433   0.335 0.737989
```

## Interaction Terms

It is easy to include interaction terms in a linear model using the `lm()` function. The syntax `lstat:black` tells R to include an interaction term between `lstat` and `black`. The syntax `lstat*age` simultaneously includes `lstat`, `age`, and the interaction term `lstat×age` as predictors; it is a shorthand for `lstat+age+lstat:age`.

NOTE: `lm.fit6=lm(medv~lstat *age ,data=Boston )`= `lm (medv~lstat+age+lstat:age ,data=Boston )`

```
lm.fit6 = lm(medv~lstat *age ,data=Boston )
summary(lm.fit6)
```

```
##
## Call:
## lm(formula = medv ~ lstat * age, data = Boston)
##
## Residuals:
```

##	Min	1Q	Median	3Q	Max
##	-15.806	-4.045	-1.333	2.085	27.552

## Interaction Terms with no code

```
lm.fit6 = lm(medv~lstat *age ,data=Boston )  
summary(lm.fit6)
```

```
##
```

```
## Call:
```

```
## lm(formula = medv ~ lstat * age, data = Boston)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max  
## -15.806  -4.045  -1.333   2.085  27.552
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept) 36.0885359  1.4698355  24.553  < 2e-16 ***  
## lstat      -1.3921168  0.1674555  -8.313  8.78e-16 ***  
## age        -0.0007209  0.0198792  -0.036  0.9711  
## lstat:age    0.0041560  0.0018518   2.244  0.0252 *  
## ---
```

## Non-linear Transformations of the Predictors

The `lm()` function can also accommodate non-linear transformations of the predictors. For instance, given a predictor  $X$ , we can create a predictor  $X^2$  using `I(X^2)`. The function `I()` is needed since the  $^$  has a special meaning in a formula; We now perform a regression of `medv` onto `lstat` and `lstat2`.

```
lm.fit7=lm(medv~lstat +I(lstat ^2))  
summary (lm.fit7)
```

```
##
```

```
## Call:
```

```
## lm(formula = medv ~ lstat + I(lstat^2))
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max  
## -15.2834  -3.8313  -0.5295   2.3095  25.4148
```

```
##
```

```
## Coefficients:
```

```
##      (Intercept)      lstat      lstat^2
```



# Non-linear Transformations of the Predictors without code

```
##  
## Call:  
## lm(formula = medv ~ lstat + I(lstat^2))  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -15.2834  -3.8313  -0.5295   2.3095  25.4148   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  42.862007   0.872084   49.15   <2e-16 ***  
## lstat        -2.332821   0.123803  -18.84   <2e-16 ***  
## I(lstat^2)    0.043547   0.003745   11.63   <2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 5.524 on 503 degrees of freedom  
## Multiple R-squared:  0.6407. Adjusted R-squared:  0.6393
```

## Quadratic fit is superior to the linear fit

`anova()` function to further quantify the extent to which the quadratic fit is superior to the linear fit. The near-zero p-value associated with the quadratic term suggests that it leads to an improved model.

```
lm.fit = lm(medv ~ lstat)
anova(lm.fit, lm.fit7)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Model 1: medv ~ lstat
```

```
## Model 2: medv ~ lstat + I(lstat^2)
```

```
##   Res.Df    RSS Df Sum of Sq      F    Pr(>F)
```

```
## 1      504 19472
```

```
## 2      503 15347  1    4125.1 135.2 < 2.2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
```

## Outcome





Here Model 1 represents the linear submodel containing only one predictor, `lstat`, while Model 2 corresponds to the larger quadratic model that has two predictors, `lstat` and `lstat2`. The `anova()` function performs a hypothesis test comparing the two models. The null hypothesis is that the two models fit the data equally well, and the alternative hypothesis is that the full model is superior. Here the F-statistic is 135 and the associated p-value is virtually zero. This provides very clear evidence that the model containing the predictors `lstat` and `lstat2` is far superior to the model that only contains the predictor `lstat`.

This is not surprising, since earlier we saw evidence for non-linearity in the relationship between `medv` and `lstat`.

# Residual plots of the quadratic model

We see that when the `lstat2` term is included in the model, there is little discernible pattern in the residuals.

```
par(mfrow=c(2,2))  
plot(lm.fit7)
```



## Higher Order Polynomials

In order to create a cubic fit, we can include a predictor of the form  $l(X^3)$ . A better approach involves using the `poly()` function to create the polynomial within `lm()`. For example, the following command produces a fifth-order polynomial fit:

```
lm.fit9=lm(medv~poly(lstat ,5))  
summary(lm.fit9)
```

```
##
```

```
## Call:
```

```
## lm(formula = medv ~ poly(lstat, 5))
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max  
## -13.5433  -3.1039  -0.7052   2.0844  27.1153
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
```

# Results

- ▶ This suggests that including additional polynomial terms, up to fifth order, leads to an improvement in the model fit!
- ▶ However, further investigation of the data reveals that no polynomial terms beyond fifth order have significant p-values in a regression fit.
- ▶ We are in no way restricted to using polynomial transformations of the predictors.

Here we try a log transformation.

```
summary (lm(medv~log(rm),data=Boston ))
```

```
##
```

```
## Call:
```

```
## lm(formula = medv ~ log(rm), data = Boston)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -19.487  -2.875  -0.104   2.837   39.816
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -76.488      5.028  -15.21  <2e-16 ***
## log(rm)       54.055      2.739   19.73  <2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
```

```
##
```

```
## Residual standard error: 6.915 on 504 degrees of freedom
```

## Qualitative Predictors

We will now examine the Carseats data, which is part of the ISLR library. We will attempt to predict Sales (child car seat sales) in 400 locations based on a number of predictors.

```
library(ISLR)
attach(Carseats )
names(Carseats )
```

```
## [1] "Sales"          "CompPrice"      "Income"         "Advertis"
## [6] "Price"          "ShelveLoc"     "Age"            "Education"
## [11] "US"
```

```
dim(Carseats)
```

```
## [1] 400 11
```



## Check variable types

```
sapply(Carseats,class)
```

```
##      Sales  CompPrice      Income Advertising  Populat  
## "numeric" "numeric"  "numeric"  "numeric"  "numer  
## ShelfLoc      Age Education      Urban  
## "factor"  "numeric" "numeric"  "factor"  "fact
```

## Qualitative Predictors

The Carseats data includes qualitative predictors such as Shelveloc, an indicator of the quality of the shelving location—that is, the space within a store in which the car seat is displayed—at each location.

The predictor Shelveloc takes on three possible values, Bad, Medium, and Good.

Given a qualitative variable such as Shelveloc, R generates dummy variables automatically. Below we fit a multiple regression model that includes some interaction terms.

```
lm.fit10 =lm(Sales~.+ Income:Advertising +Price:Age ,data=C  
summary (lm.fit10)
```

```
##
```

```
## Call:
```

```
## lm(formula = Sales ~ . + Income:Advertising + Price:Age,
```

```
##
```

```
## Residuals:
```

## Output without code

```
##  
## Call:  
## lm(formula = Sales ~ . + Income:Advertising + Price:Age,  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -2.9208 -0.7503  0.0177  0.6754  3.3413   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)    6.5755654   1.0087470   6.519 2.22e-11   
## CompPrice      0.0929371   0.0041183  22.567 < 2e-16   
## Income         0.0108940   0.0026044   4.183 3.57e-05   
## Advertising    0.0702462   0.0226091   3.107 0.00203   
## Population     0.0001592   0.0003679   0.433 0.66533   
## Price         -0.1008064   0.0074399 -13.549 < 2e-16   
## ShelveLocGood  4.8486762   0.1528378  31.724 < 2e-16   
## ShelveLocMedium 1.9532620   0.1257682  15.531 < 2e-16
```

## The contrasts() function

returns the coding that R uses for the dummy variables.

```
## The following objects are masked from Carseats (pos = 3)
```

```
##
```

```
##      Advertising, Age, CompPrice, Education, Income, Popu
```

```
##      Price, Sales, ShelveLoc, Urban, US
```

```
##           Good Medium
```

```
## Bad           0       0
```

```
## Good          1       0
```

```
## Medium        0       1
```

NOTE: Use ?contrasts to learn about other contrasts, and how to set them.

- ▶ R has created a ShelfLocGood dummy variable that takes on a value of 1 if the shelving location is good, and 0 otherwise.
- ▶ It has also created a ShelfLocMedium dummy variable that equals 1 if the shelving location is medium, and 0 otherwise.
- ▶ A bad shelving location corresponds to a zero for each of the two dummy variables.
- ▶ The fact that the coefficient for ShelfLocGood in the regression output is positive indicates that a good shelving location is associated with high sales (relative to a bad location).
- ▶ And ShelfLocMedium has a smaller positive coefficient, indicating that a medium shelving location leads to higher sales than a bad shelving location but lower sales than a good shelving location.

# Writing Functions

We now create the function. Note that the `+` symbols are printed by R and should not be typed in. The `{` symbol informs R that multiple commands are about to be input. Hitting Enter after typing `{` will cause R to print the `+` symbol. We can then input as many commands as we wish, hitting Enter after each one. Finally the `}` symbol informs R that no further commands will be entered.

- Craete a function called LoadLibraries:

```
LoadLibraries=function (){  
  library(ISLR)  
  library(MASS)  
  print("The libraries have been loaded")  
}
```

Now if we type in LoadLibraries, R will tell us what is in the function.

```
## function (){  
##   library(ISLR)  
##   library(MASS)  
##   print("The libraries have been loaded")  
## }
```

If we call the function, the libraries are loaded in and the print statement is output.

```
## [1] "The libraries have been loaded"
```